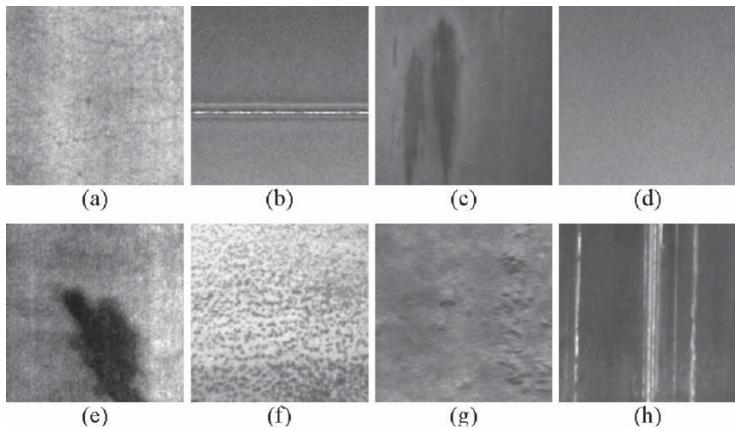
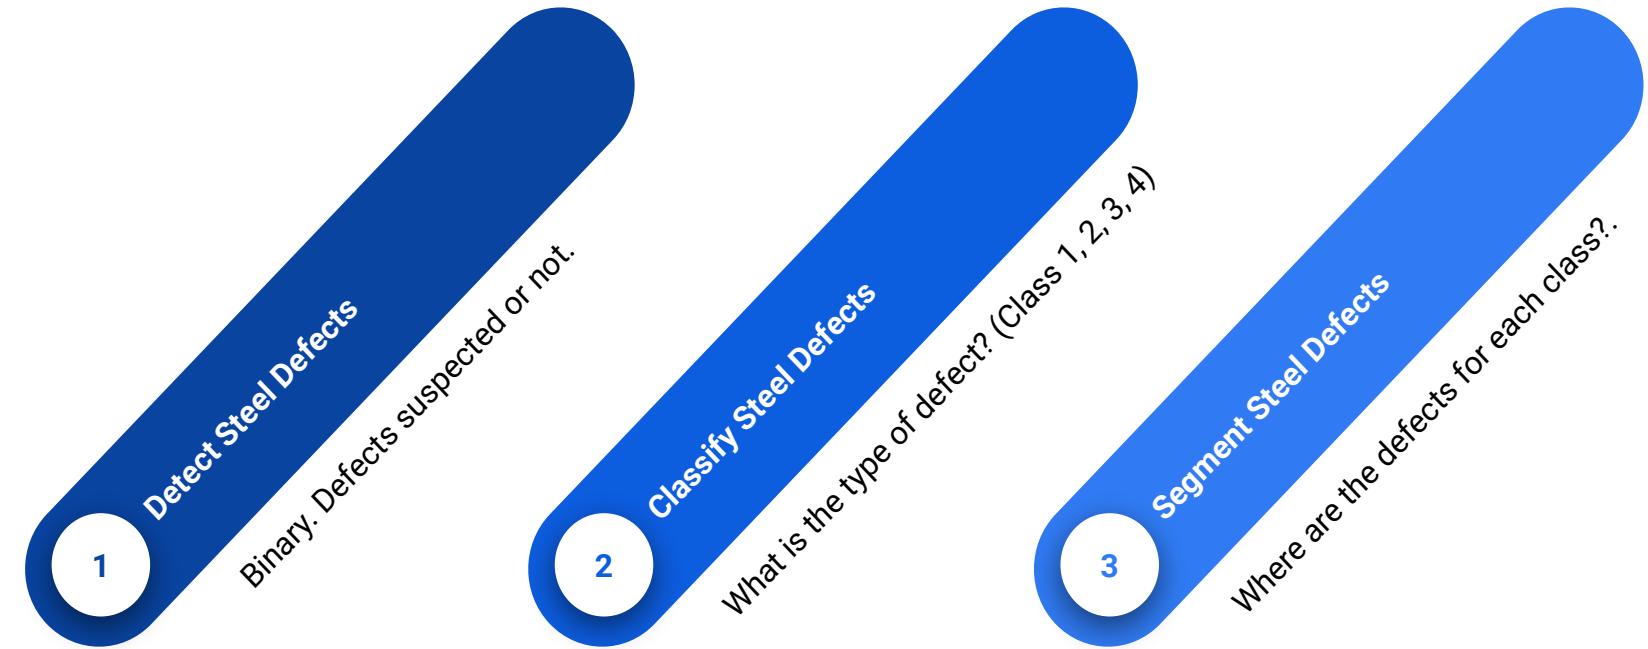


# Detection, Classification, and Segmentation of Steel Defects

---

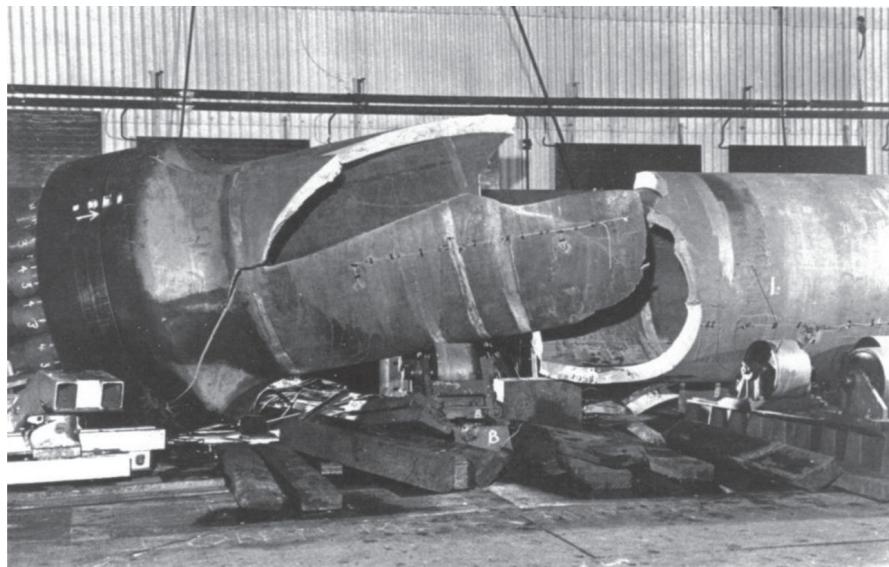
Farhan Wadia, Hiranya Maharaja,  
Wenju Zhang, Zhipu Zhang

# Introduction & Problem





# Significance



# Data

<u>ImageId</u>	<u>ClassId</u>	<u>EncodedPixels</u>
----------------	----------------	----------------------

0007a71bf.jpg_1		
0007a71bf.jpg_2		
0007a71bf.jpg_3		18661 28 18863 82 19091 110 19347 110 19
0007a71bf.jpg_4		
000a4bcdd.jpg_1		37607 3 37858 8 38108 14 38359 20 38610
000a4bcdd.jpg_2		
000a4bcdd.jpg_3		
000a4bcdd.jpg_4		
000f6bf48.jpg_1		
000f6bf48.jpg_2		
000f6bf48.jpg_3		
000f6bf48.jpg_4		131973 1 132228 4 132483 6 132738 8 1329
0014fce06.jpg_1		
0014fce06.jpg_2		
0014fce06.jpg_3		229501 11 229741 33 229981 55 230221 77
0014fce06.jpg_4		
001982b08.jpg_1		
001982b08.jpg_2		
001982b08.jpg_3		
001982b08.jpg_4		



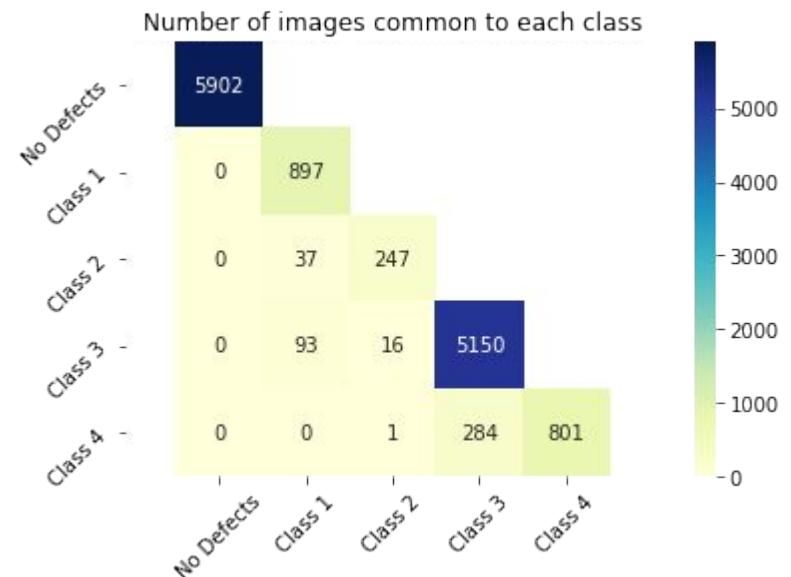
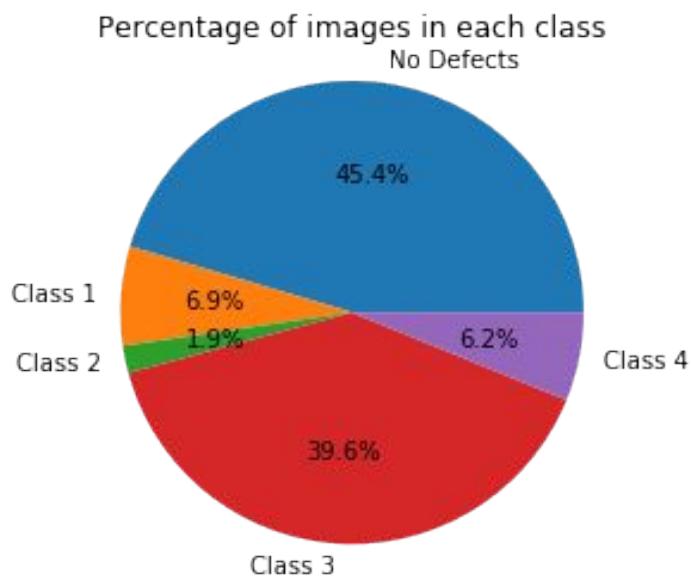
1	7	13	19	25	31
2	8	14	20	26	32
3	9	15	21	27	33
4	10	16	22	28	34
5	11	17	23	29	35
6	12	18	24	30	36

12568 images (256 x 1600) and  $12568 \times 4$  run-length encodings.

**EncodedPixels:** 8 3 18 4 33 2

# Data

- Classes 1-4 split 70-15-15 for stage 2 training, validation, and testing
  - 4760 training images, 1055 validation images, and 1055 test images



# Data Processing

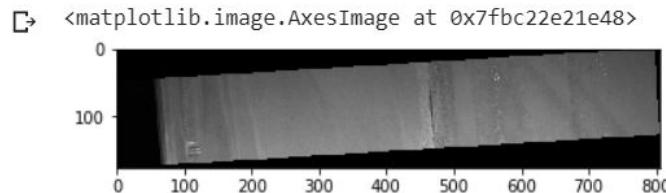
- Convert run-length encodings to one-hot masks
- Downsize images (trained with 96 x 576)
- Randomly apply (25% of the time):
  - Horizontal flips
  - Vertical flips
  - Rotations
- Convert to tensor

**Issue:** How to apply a random transformation to an image, but also apply that same transformation to its segmentation mask?

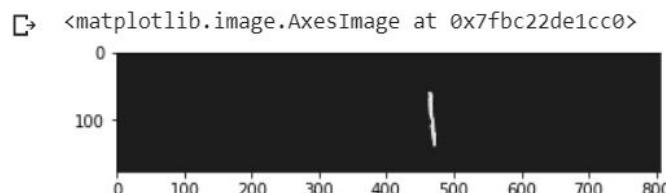
**Solution:** Create a custom dataset class to load and process images and masks before passing to the data loader



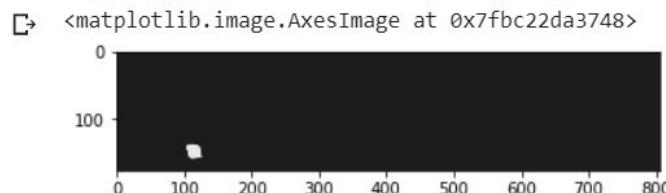
```
[23] # a known image with defects in class 1 and 3
image, masks = train_dataset[4137]
plt.imshow(image.numpy(), cmap = 'gray')
```



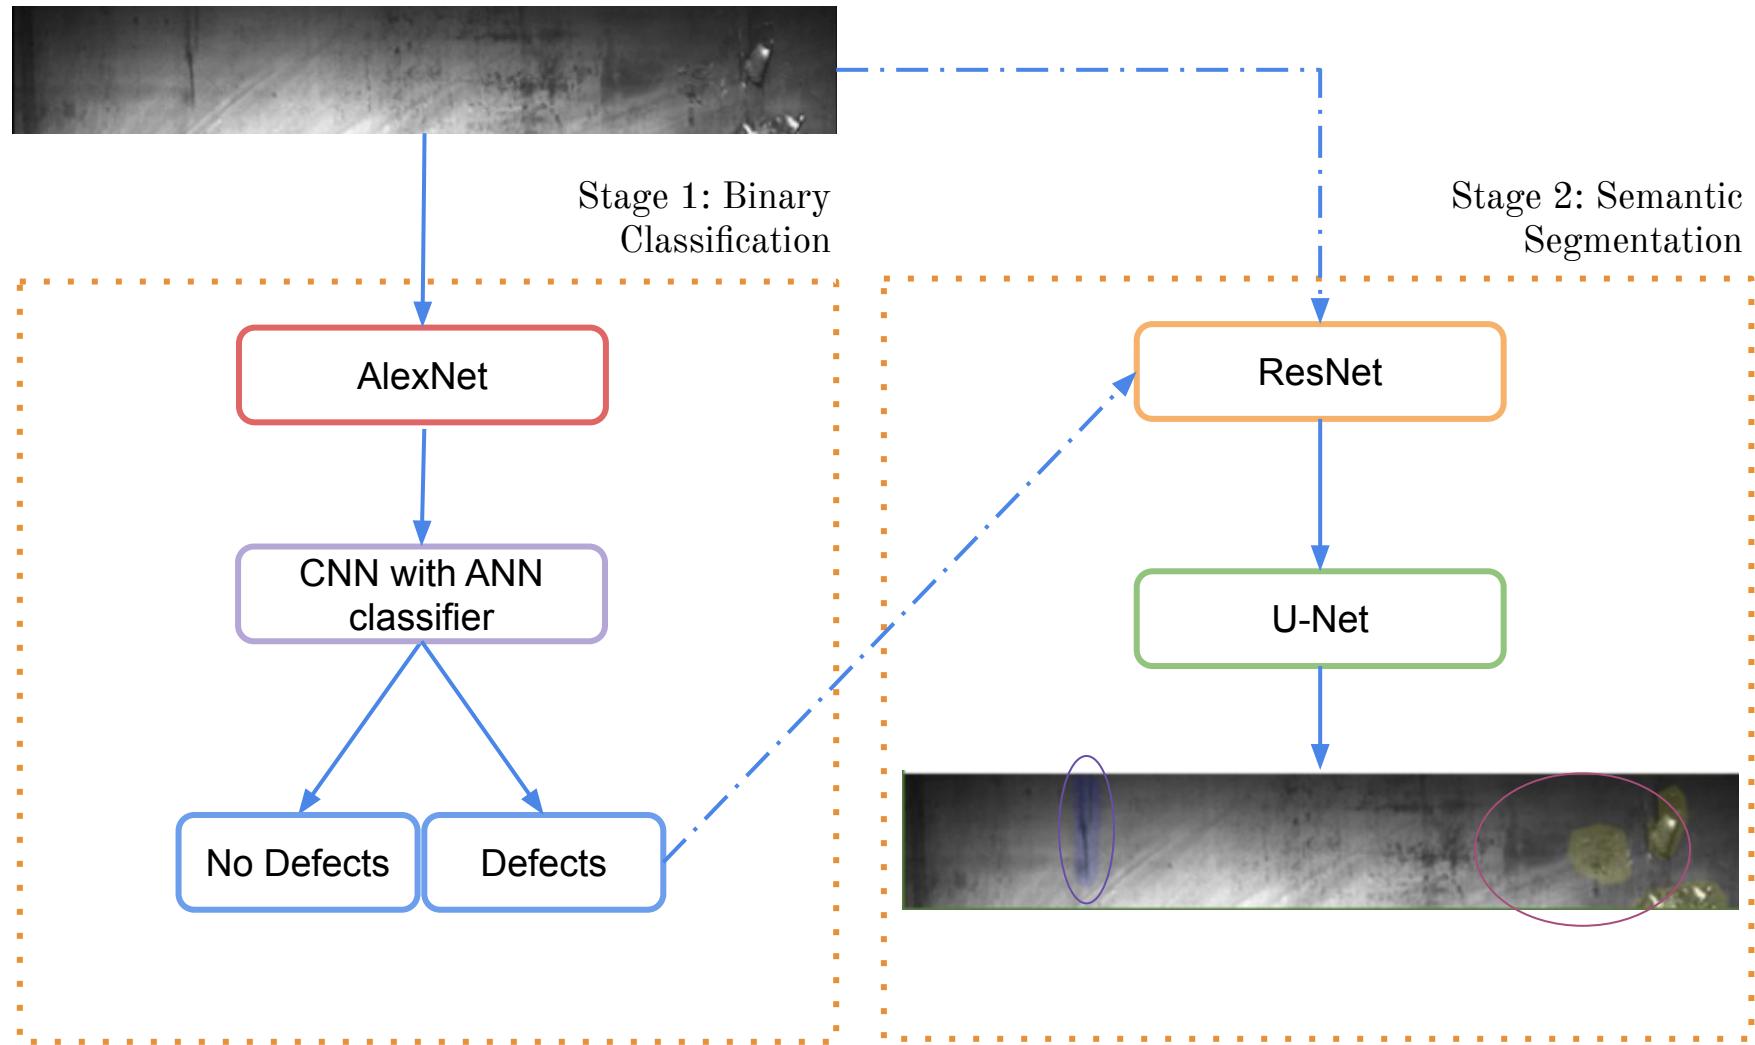
```
[24] #verify that class 3 map is transformed correctly
plt.imshow(masks.numpy()[2], cmap = 'gray')
```



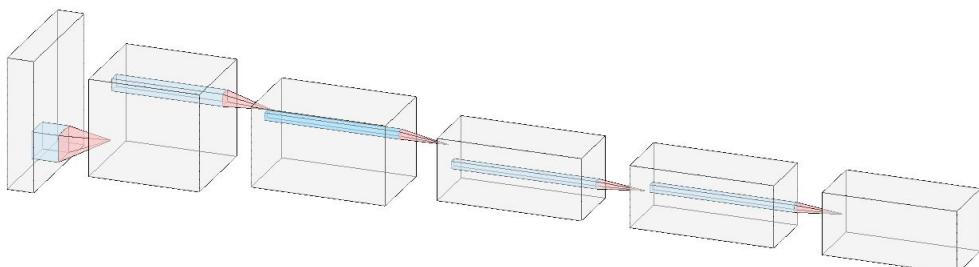
```
[25] #verify that class 1 map is transformed correctly
plt.imshow(masks.numpy()[0], cmap = 'gray')
```



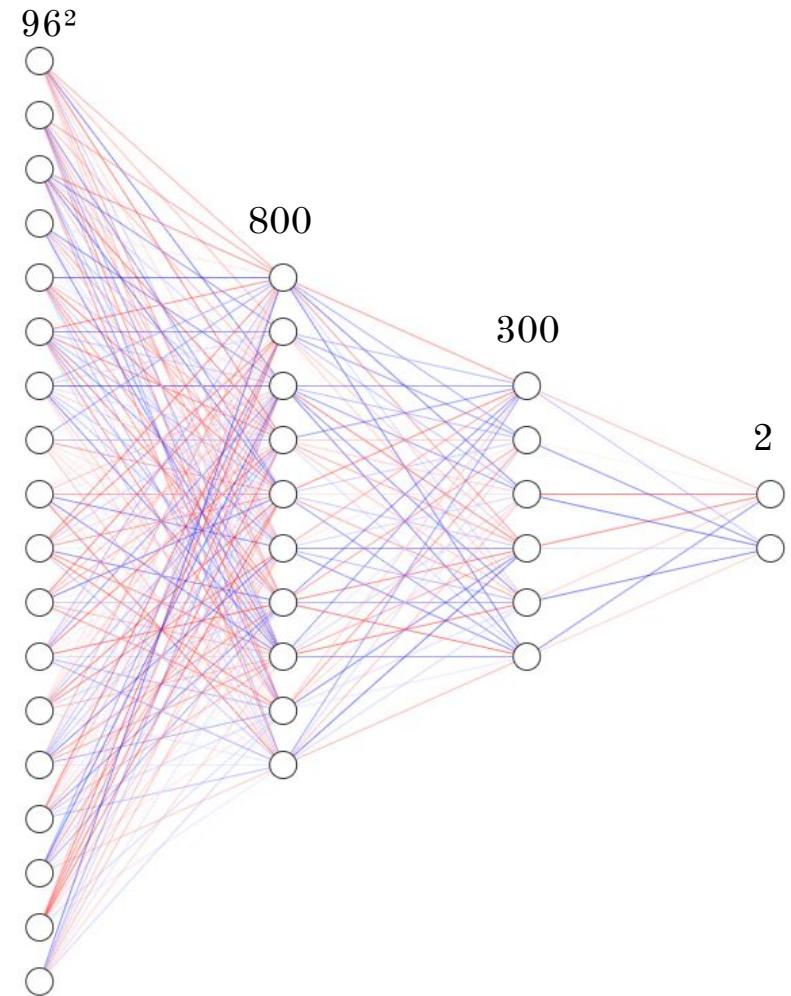
# Models Overview



# Stage 1: Binary Classification



Pre-trained AlexNet feature extractor

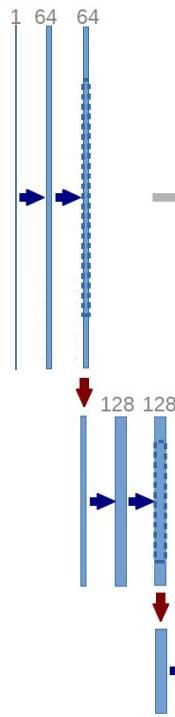


ANN Classifier

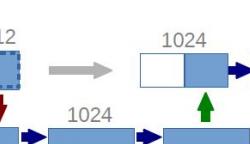
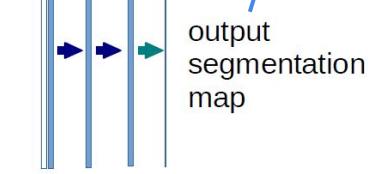
# Stage 2: ResNet U-Net Semantic Segmentation

Batch Size x 3 x 96 x 576 Images

Compute pre-trained ResNet features



Batch Size x 4 x 96 x 576 Masks



- conv 3x3, ReLU
- copy and crop
- ↓ max pool 2x2
- ↑ up-conv 2x2
- conv 1x1

# Stage 2: Hyperparameter Tuning

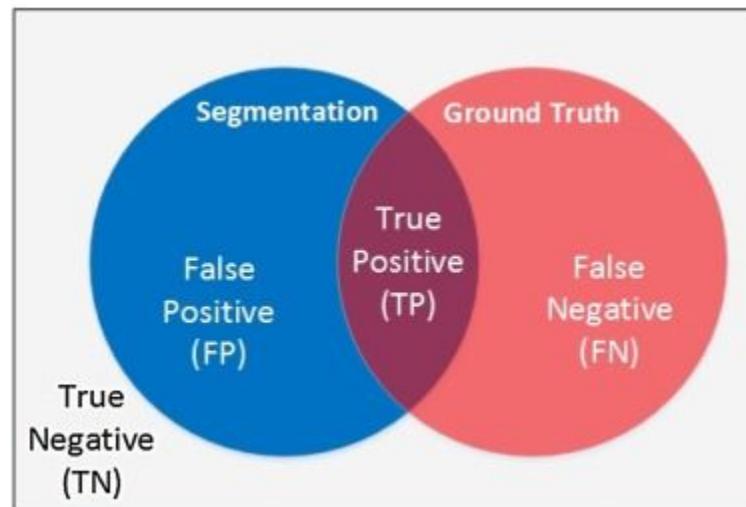
- Tried U-Net model without pre-trained ResNet features
- Changing number of epochs
- Changing learning rate
- Changing loss function
  - Weighted combination of BCE and Dice
  - Best model used only BCE



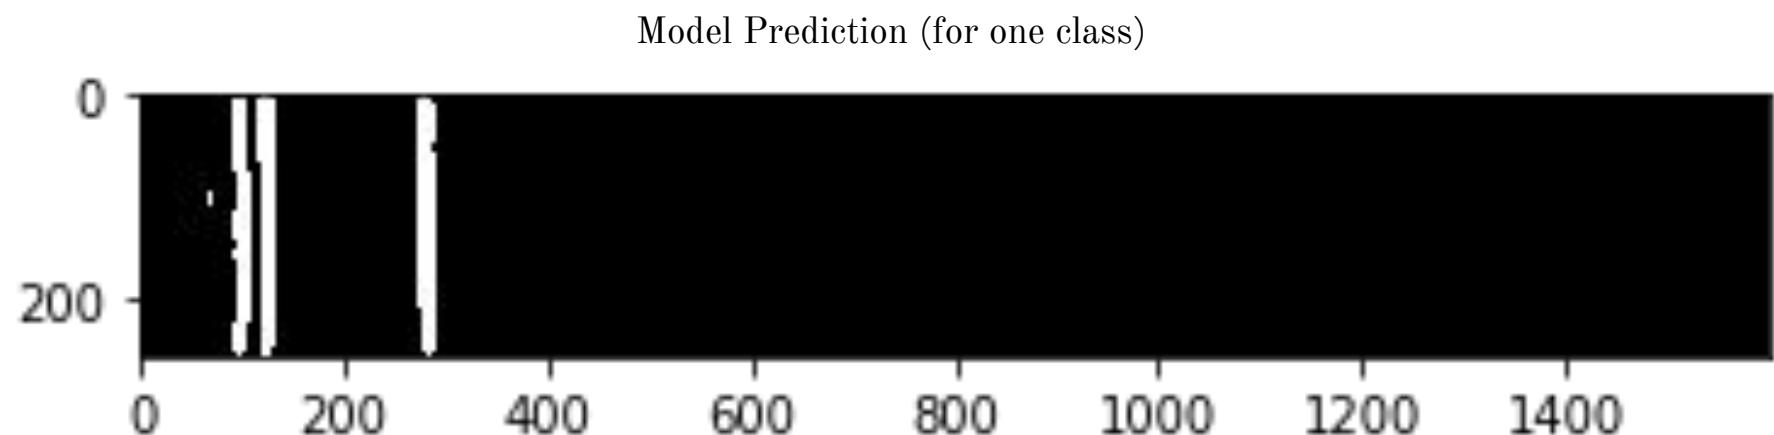
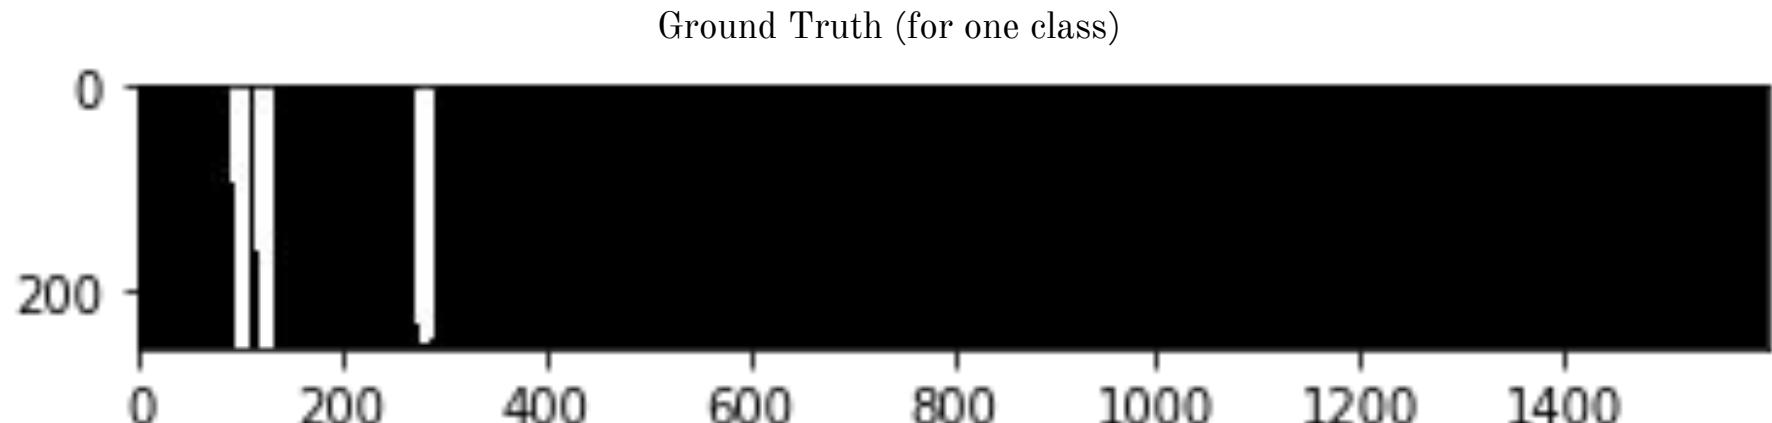
# Quantifying Accuracy: Dice Coefficient

- Measures the accuracy of the semantic segmentation model
- Indicates how similar the predictions are to the targets pixel-by-pixel within each class
- Average the dice coefficients over all the classes and the number of images

$$\frac{2 \cdot (\text{number of true positives})}{\text{number of true positives} + \text{number of false positives} + \text{number of false negatives}}$$



# So what does a result look like?



# Test Results

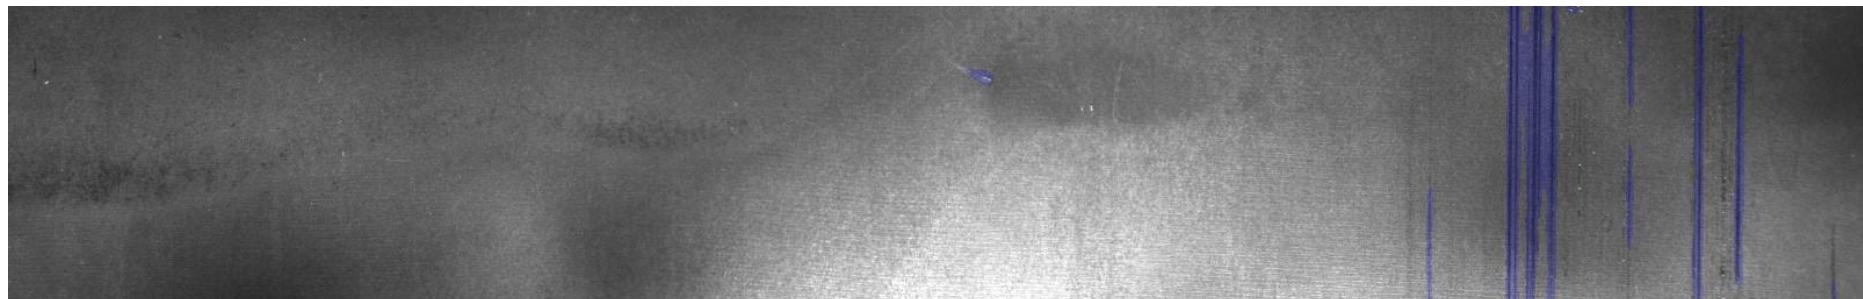
	SVM Baseline	Model
Stage 1: Binary Detection Accuracy	91%	100%
Stage 2: Semantic Segmentation Accuracy	72%	83% (Mean Dice)

# Examples from Test Results

Ground Truth



Predicted



**Mean Dice Coefficient: 92.4%**

# Examples from Test Results

Ground Truth



Predicted



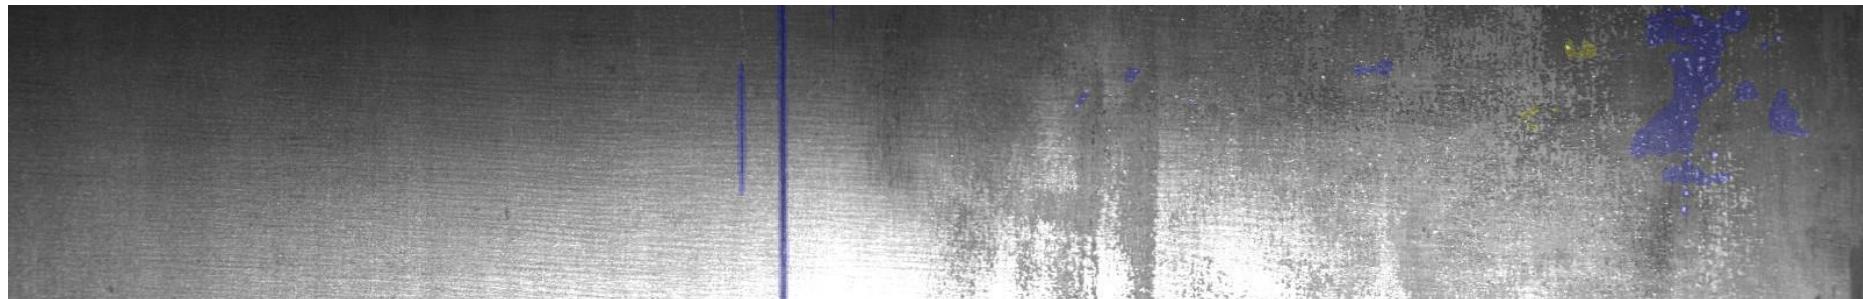
**Mean Dice Coefficient: 87.3%**

# Examples from Test Results

Ground Truth



Predicted



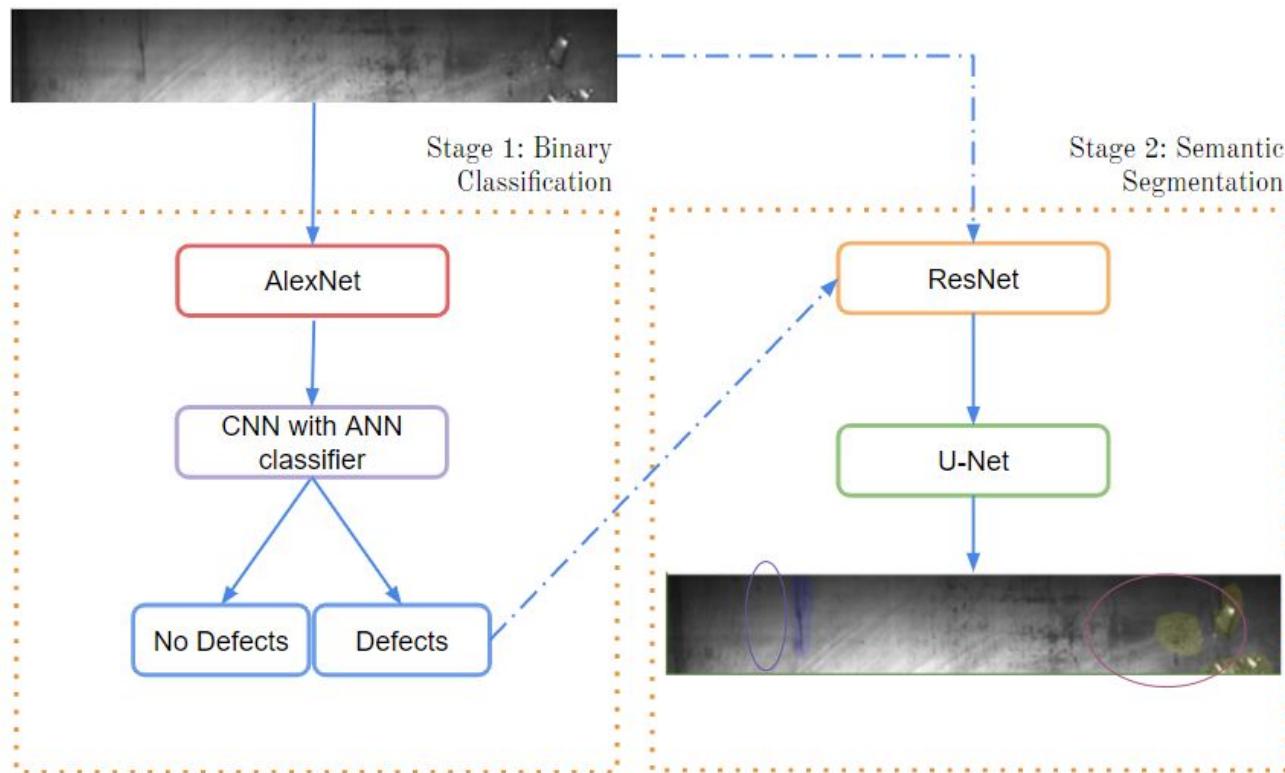
**Mean Dice Coefficient: 56.3%**

# Qualitative Results - Reasons for Success

- Usage of transfer learning:
  - AlexNet
  - ResNet 18
- ResNet skip connections allow for deeper learning without vanishing gradients
- U-Net is well-known for semantic segmentation in the biomedical industry

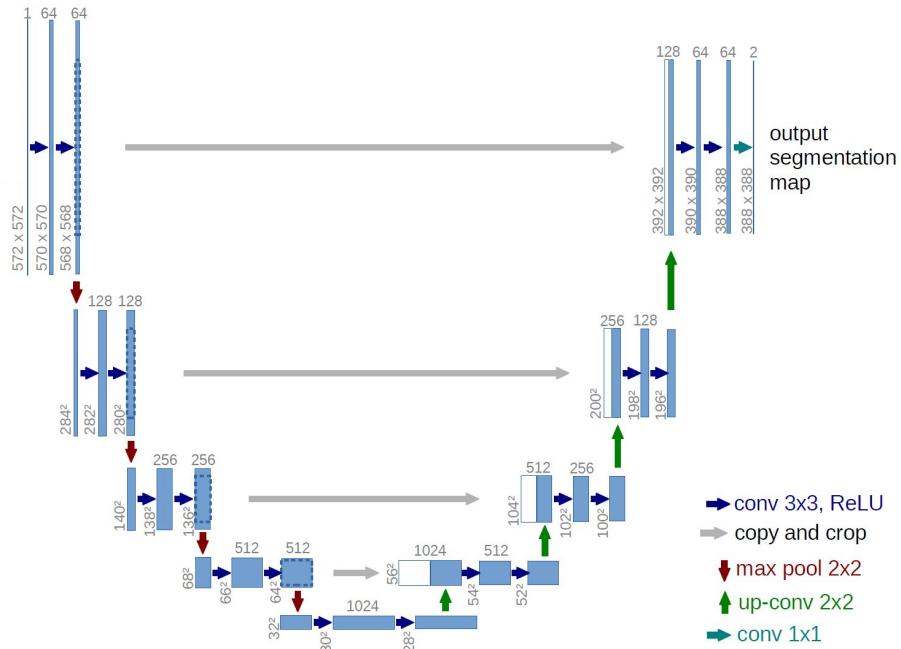
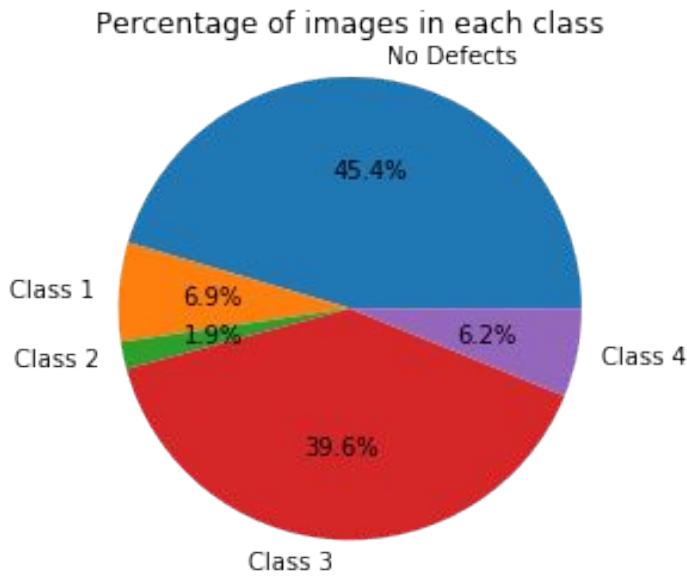
# Qualitative Results - Comparisons

- No significant difference in accuracy for model with Stage 1 vs. without
  - ResNet - Enough layers to go from 4 to 5 classes
  - Stage 2 convolutions also extract features - similar to AlexNet
- Sanity check - both models should result in the same answer



# Qualitative Results - Reasons for Error

- Dataset is imbalanced
- Information lost when encoding/decoding features through U-Net



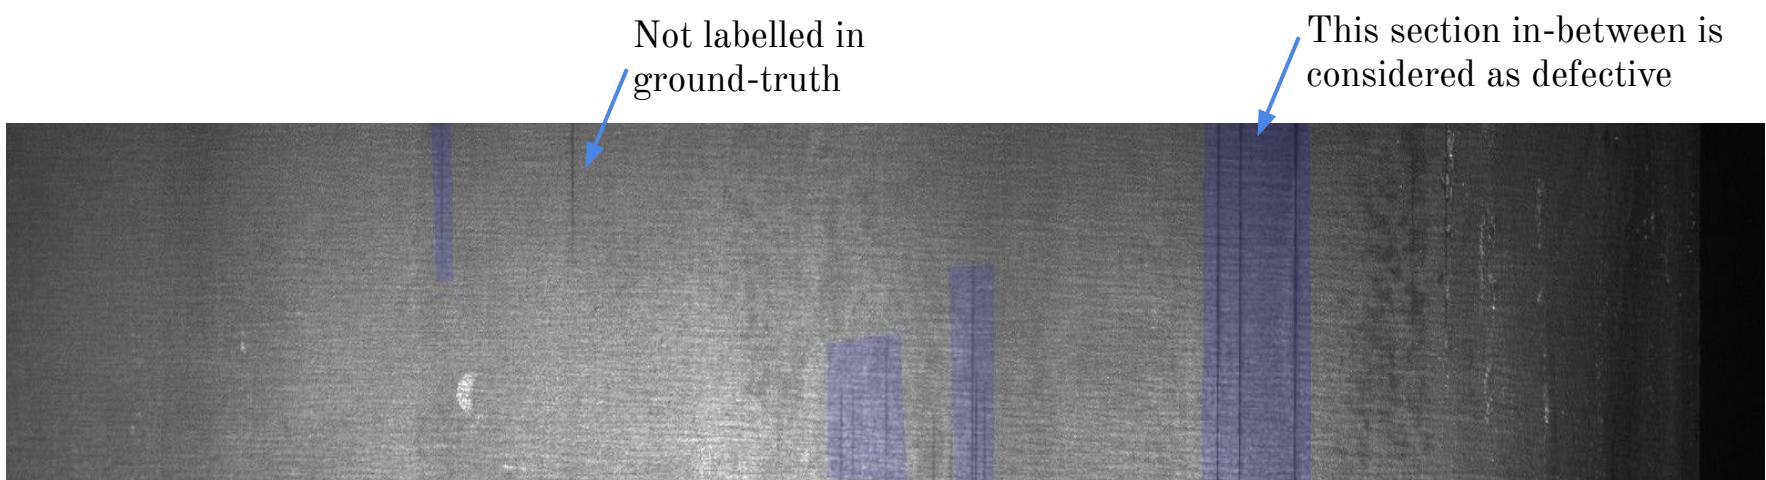
# Discussion / Insights

- Misleading Labels
  - The image below is considered to have no defects because this type of defect is not part of class 1, 2, 3, or 4
- Incomplete Number of Classes
  - Steel defects cannot be segmented into only 4 classes; there are many more types of defects than the ones our model was trained to identify



# Discussion / Insights

- Mislabeling
  - Datasets were manually labeled - ground truth labels might be inaccurate
- Manufacturer Bias
  - All data is of new steel sheets from one manufacturer
  - Potentially cannot be extrapolated to samples from other manufacturers
  - Will not extrapolate to more complex in-field situations (e.g. rail crack detection, rebar crack detection etc.)
- Overfitting Potential
  - Images that appear grayscale are being passed in as RGB



# Demonstration

(Kaggle-provided test set  
of 1800 images, ground  
truth unknown)

